



# Deep Neural Networks for Encrypted Inference with TFHE

Andrei Stoian<sup>(✉)</sup>, Jordan Frery, Roman Bredehoft, Luis Montero,  
Celia Kherfallah, and Benoit Chevallier-Mames

Zama, Paris, France  
[hello@zama.ai](mailto:hello@zama.ai)  
<http://zama.ai>

**Abstract.** Fully homomorphic encryption (FHE) is an encryption method that allows to perform computation on encrypted data, without decryption. FHE preserves the privacy of the users of online services that handle sensitive data, such as health data, biometrics, credit scores and other personal information. A common way to provide a valuable service on such data is through machine learning and, at this time, neural networks are the dominant machine learning model for unstructured data.

In this work we show how to construct Deep Neural Networks (DNN) that are compatible with the constraints of TFHE, an FHE scheme that allows arbitrary depth computation circuits. We discuss the constraints and show the architecture of DNNs for two computer vision tasks. We benchmark the architectures using the Concrete stack (<https://github.com/zama-ai/concrete-ml>), an open-source implementation of TFHE.

## CSCML 2023 Submission

This paper is submitted as a **short paper**.

## 1 Introduction

Neural Networks (NNs) are machine learning (ML) models that have driven the recent expansion of the field of Artificial Intelligence (AI). Their performance on unstructured data such as images, sound and text is unmatched by other ML techniques [19]. Though training deep NNs is more difficult, they obviate the need for complex feature engineering and process raw data directly, making them easier to deploy in production. Applications of NNs include image classification, face recognition, voice assistants, and search engines, tools which today are a staple of the user experience online. Deployment of such models in SaaS applications raises a security risk: they are a target of malevolent entities that seek to steal the sensitive user data these models process.

Privacy-preserving technologies, such as multi-party computing (MPC) and fully homomorphic encryption (FHE), provide a solution to the risk of data leaks, eliminating it by design. Notably, FHE encrypts user data and allows a

third party to process the data in its encrypted form, without needing to decrypt it. Only the owner of the secret-key can decrypt the result of the computation. Thus, an attacker can only steal encrypted data they can not decrypt.

In this work we show how to build NNs that are FHE compatible, while minimizing the cryptography knowledge needed by the machine learning practitioner. We based our work on the Concrete Library [9] which uses TFHE [8], works over integers, provides a fast *programmable* bootstrapping (PBS) mechanism, and performs exact computation.

## 2 Related Work

Several alternative approaches exist for neural network inference over encrypted data. Many of these works rely on encryption schemes such as CKKS [7] and YASHE [4] which support both floating point and integer computation. Most works on NNs use quantized models where the weights and activations are fixed point or integers. With the exception of [15], these works avoid bootstrapping in order to reduce inference time. Avoiding bootstrapping has the disadvantage of limiting the depth of the models that can be computed in FHE.

CryptoNets [12] uses YASHE which supports the computation of polynomials of encrypted values. CryptoNets are NNs quantized to integers (of 5–10 bits) with activation functions expressed as low-degree polynomials. CryptoNets achieve 99% accuracy on MNIST using a three layer network with an inference time of 570 s/image. To speed-up this technique, [6] optimize the ciphertext representation of clear values to leverage packing techniques specific to CKKS.

FHE-DiNN [5], a TFHE based approach, quantizes inputs, intermediate values and weights to binary values. During training the activation functions are set to `hardSigmoid`, while in inference they are swapped for the `sign` function. However, binary NNs are hard to train and do not perform well in many ML tasks such as object detection and speech processing.

Another TFHE approach, SHE [16], uses bit series representation of encrypted values and boolean gates. They run NNs that fit within a maximum multiplicative depth budget and, by avoiding expensive multi-bit PBSs, they achieve inference of a ShuffleNet on ImageNet with a latency of 18,000 s/image. They rely on logarithmic quantization of weights which allows to reduce multiplicative depth for the convolution layers by using bit-shifts. Sums, `relu` and `maxpool` are computed using boolean gates. Another TFHE-based approach, combining both boolean and integer ciphertext representations, [11], was shown to take advantage of the strengths of both types of representation.

Leveled approaches such as SHE and CryptoNets are limited by the maximum multiplicative depth budget, which, in turn, limits the supported network types and their depth. Moreover, to speed up the computations, many works that use CKKS set crypto-system parameters in a way that makes the computations approximate in the sense that noise corrupts some of the message bits.

In this work we propose an approach to train arbitrary NNs which can have any depth, number of neurons and activation functions. Furthermore, our approach performs exact computation in FHE: the noise of the encryption scheme

does not corrupt the values that are processed. Thus results in FHE are the same as in the clear - there is no degradation of accuracy when moving to encrypted inference - which is a major advantage when putting models in production.

### 3 Neural Network Training for Encrypted Inference

Training NNs is usually done in floating point, but most FHE schemes, including TFHE, only support integers. Consequently, quantization must be used, and two main approaches exist:

1. Post-training quantization is commonly used [12, 16], but, in this mode, NNs lose accuracy when the quantization bit-width is lower than 7–8 bits. With per-channel quantization, or logarithmic quantization, which are more complex to implement, as few as 4 bits were used for weights and activations without loss of accuracy [20].
2. Quantization-aware Training (QAT), used in this work and in [5], is an approach that adds quantizers to network activations and weights during training. QAT enables extreme quantization with less than 4 bit weights and activations.

To support arbitrarily deep NNs and any activation function, we make use of the PBS [10] mechanism of TFHE. PBS reduces the noise in accumulators of ciphertext leveled operations (addition, multiplication with clear constants) but also allows to apply a lookup-table (TLU) on its input ciphertext.

The TFHE PBS mechanism has a rather high computational cost, and this cost depends on the number of bits of the encrypted value to be bootstrapped. It is convenient to keep the accumulator size low, in order to speed up the PBS computation. However, reducing accumulator bit-width has a negative impact on network prediction performance, so a compromise needs to be found.


We describe here a QAT strategy that can process all the intermediate encrypted values as integers. In this way, training an FHE compatible network becomes purely a machine learning problem and no cryptography knowledge is needed by the practitioner. To build a TFHE compatible NN, the constraints on the network architecture are the following:

- All layers that sum or multiply two encrypted values, such as convolution **conv** and fully-connected **fc**, must have quantized inputs. This is easily achieved using QAT frameworks.
- The bit-width of the accumulators of layers such as **conv**, **fc** must be bounded. To achieve this, similarly to [1], we use pruning of neuron synapses.

To control the accumulator bit-width while keeping the training dynamics stable, we use  $L^1$ -norm unstructured pruning. Figure 1 shows the impact of pruning on the accumulator size for two quantization modes: narrow and wide range.

While the inputs of **conv** and **fc** layers need to be quantized, it is possible to use floating point layers for all univariate operations such as batch normalization, quantization, and activations.

In our FHE compatible NNs the outputs of a `conv` or `fc` are processed by a sequence of univariate operations that ends with quantization. This sequence of functions takes integers and has integer outputs, but the intermediary computations in these operations can use float parameters. Thus, batch normalization, activation functions, neuron biases and any other univariate transformation of `conv` or `fc` outputs does not need quantization. Figure 1 shows the architecture of the network during training and inference.



**Fig. 1.** Left: accumulator size while varying the number of active neurons during pruning for a 3-layer fully-connected network with 2 bit weights and activations. Two quantization modes are shown: Narrow range uses values  $[-2^{b-1} + 1, 2^{b-1} - 1]$ , while Wide range uses  $[-2^{b-1}, 2^{b-1} - 1]$ . Right: the structure of a 2 layer convolutional network in training and during inference. Univariate layers are fused to table-lookups, implemented with PBS.

## 4 Neural Network Inference Using TFHE

Inference of our FHE NNs is based on quantized implementations of NN operators that add or multiply together encrypted values. Convolutional, fully connected and average pooling layers use the quantized formulation from [13]. Since uniform quantization is used, we can define a quantized value  $r$  as  $r = S(q - Z)$  where  $S$  is the quantization scale,  $Z$  is the quantization zero-point and  $q$  is the integer representation of the value. Next, the fully connected layer, with inputs  $x$ , weights  $w$  and outputs  $o$ , with per-tensor quantization parameters  $(S_x, Z_x)$ ,  $(S_w, Z_w)$  can be written as:

$$S_o(q_o^k - Z_o) = \sum_{i=0}^N S_x(q_x^i - Z_x) S_w(q_w^{(i,k)} - Z_w) + b^k \quad (1)$$

where  $k$  is the index of a neuron in the layer and  $N$  is the number of connections of the neuron and  $b^k$  is the bias of the  $k$ -th neuron. A convolutional layer can be

expressed by extending the sum to the height, width and channel dimensions. Equation 2 can be re-written to separate integer and floating point computations (note that zero-points  $Z_x, Z_o, Z_w$  are integers).

$$q_o^k = b^k + Z_o + \frac{S_x S_w}{S_o} \sum_{i=0}^N (q_x^i - Z_x)(q_w^{(i,k)} - Z_w) \quad (2)$$

Therefore we can separate the equation in a floating point univariate function  $f$  and a sum over products of encrypted inputs and clear weights:

$$q_o^k = f(\Sigma) \text{ where } f(q) = b^k + Z_o + \frac{S_x S_w}{S_o} q \text{ and } \Sigma = \sum_{i=0}^N (q_x^i - Z_x)(q_w^{(i,k)} - Z_w) \quad (3)$$

The univariate function  $f$  in Eq. 3 takes integer inputs. We compose this function with the batch-normalization, and, finally, with the quantization function  $Q(x) = \text{floor}\left(\frac{x}{S_x}\right) + Z_x$ . Thus  $f$  becomes a function defined on  $\mathbb{Z}$ , with values in  $\mathbb{Z}$  and can be implemented as a lookup table with a PBS in FHE, without any loss of precision. Our approach - fusing all univariate functions - contrasts with other works [11, 13] which, instead, prefer to express the batch normalization and quantization with integer operations. However, expressing all these univariate operations with integers only adds to the difficulty of the quantization process and decreases accuracy. Our PBS approach allows fusing all univariate floating point operations applied to an encrypted integer to a lookup-table, thus reducing quantization constraints.

The complete NN computation can now be expressed over integers using the following operations: multiplication of an encrypted value and a clear constant, sums of encrypted integer values, table lookup of encrypted integer values. In our implementation of TFHE, Concrete, we encode integers in two different ways: integers up to 8 bits are encoded into a single ciphertext, and integers between 9–16 bits are encoded with a CRT representation into several ciphertexts as described in [3]. This contrasts to previous works, such as [16], that encode each bit of an integer as an individual ciphertext and use boolean gates to build arithmetic circuits.

An automated optimization process [3] determines the cryptographic parameters of the circuit, based on several factors: (1) the *circuit bit-width*, defined as the minimum bit-width necessary to encode the largest integer value obtained anywhere in the NN's integer-based evaluation, (2) the maximum 2-norm of the integer weight tensors of the layers, and (3), the desired probability of error of the PBS. The optimization process determines the cryptosystem parameters (LWE dimension, polynomial size, GLWE dimension, etc. - see [14]) to ensure a fast execution, the target probability of failure and the security level (using the lattice-estimator [2]). We set the PBS error probability sufficiently low to ensure full correctness of the results, i.e. the results in the clear are always the same as those in FHE, up to a user-defined error-rate, e.g.  $10^{-6}$ , for one full NN inference.

## 5 Experimental Results

The networks were implemented in PyTorch with Brevitas [18] and converted to FHE with Concrete-ML [17]. We ran experiments on two datasets with several neural network architectures, in two quantization modes (see Fig. 1, left). In Table 1, the latency results for encrypted inference are given. Unless otherwise indicated by a footnote, the test machine had an Intel i7-11800H CPU with 8 cores (16 threads were used). For one of the configurations, as indicated in the table, the first layer of the network was executed in the clear to keep the overall accumulator size low.

**Table 1.** Experimental results obtained with Brevitas and Concrete-ML. Unless otherwise stated, the timing is obtained on an 8-core machine.

Network	Quant. bits	Active Neurons	Narrow range	Data-set	Circuit bit-width	Accuracy	Inference time (s)
3-layer FCNN <sup>a</sup>	2/2	150	Yes	MNIST	6	92.2%	31
3-layer FCNN	2/2	90	No	MNIST	7	96.5%	77
3-layer FCNN	2/2	190	No	MNIST	8	97.1%	300
LeNet	2/2	190	No	MNIST	8	97.6%	2780
6-layer CNN <sup>b</sup>	2/2	190	No	MNIST	8	98.7%	5072
VGG-9 <sup>c</sup>	2/2 <sup>d</sup>	all	Yes	CIFAR10	8	62.3%	624 <sup>e</sup>
VGG-9	6/6	110	No	CIFAR10	15	90.0%	25000 <sup>e</sup>
VGG-9	6/6	110	No	CIFAR100	15	69.0%	25000 <sup>e</sup>

<sup>a</sup> Three FC layers with 192, 192 and 10 neurons

<sup>b</sup> Four conv layers with 8,8,16 and 16 filters followed by a FC layer with 120 neurons and a final FC layer for classification

<sup>c</sup> Six conv layers with: 64, 64, 128, 128, 256, 256 filters, followed by two 512 neuron FC layers and a final FC layer for classification

<sup>d</sup> Inputs are quantized in 8 bits and the first layer is executed non-encrypted. All activations use 2 bits

<sup>e</sup> On a c6i.metal AWS instance with 128 vcpu.

## 6 Conclusion

Our approach to encrypted inference for neural networks shows several advantages over other methods. First, we believe our method is easier to use than other works, since the problem of making an FHE compatible network becomes strictly an ML problem and no cryptography knowledge is needed. Second, the computations in FHE are correct with respect to the computations in the clear and, using TFHE, noise does not corrupt the encrypted values. Thus, once a network is trained incorporating the quantization constraints, the accuracy that is measured on clear data will be the same as that on encrypted data. Finally, our approach, using PBS, shows competitive accuracies in FHE and allows to

convert arbitrary-depth networks using any activation function to FHE. Networks up to 9 layers were shown, but deeper NNs can easily be implemented. Code for the MNIST, CIFAR10 and CIFAR100 classifiers is available<sup>1</sup>.

Many possible strategies can be employed to improve upon this work, in order to support larger models, such as ResNet, on larger data-sets like ImageNet. For example, a better pruning strategy could decrease the PBS count, per-channel quantization can improve accuracy, and faster step functions in FHE could improve the overall speed.

## References

1. Aharoni, E., et al.: HE-PEx: efficient machine learning under homomorphic encryption using pruning, permutation and expansion. CoRR abs/2207.03384 (2022). <https://doi.org/10.48550/arXiv.2207.03384>
2. Albrecht, M.R., Player, R., Scott, S.: On the concrete hardness of learning with errors. Cryptology ePrint Archive, Paper 2015/046 (2015). <https://eprint.iacr.org/2015/046>
3. Bergerat, L., et al.: Parameter optimization & larger precision for (T)FHE. Cryptology ePrint Archive, Paper 2022/704 (2022). <https://eprint.iacr.org/2022/704>
4. Bos, J.W., Lauter, K., Loftus, J., Naehrig, M.: Improved security for a ring-based fully homomorphic encryption scheme. In: Stam, M. (ed.) IMACC 2013. LNCS, vol. 8308, pp. 45–64. Springer, Heidelberg (2013). [https://doi.org/10.1007/978-3-642-45239-0\\_4](https://doi.org/10.1007/978-3-642-45239-0_4)
5. Bourse, F., Minelli, M., Minihold, M., Paillier, P.: Fast homomorphic evaluation of deep discretized neural networks. In: Shacham, H., Boldyreva, A. (eds.) CRYPTO 2018. LNCS, vol. 10993, pp. 483–512. Springer, Cham (2018). [https://doi.org/10.1007/978-3-319-96878-0\\_17](https://doi.org/10.1007/978-3-319-96878-0_17)
6. Brutzkus, A., Gilad-Bachrach, R., Elisha, O.: Low latency privacy preserving inference. In: Chaudhuri, K., Salakhutdinov, R. (eds.) Proceedings of the 36th International Conference on Machine Learning, ICML 2019, 9–15 June 2019, Long Beach, California, USA. Proceedings of Machine Learning Research, vol. 97, pp. 812–821. PMLR (2019). <http://proceedings.mlr.press/v97/brutzkus19a.html>
7. Cheon, J.H., Kim, A., Kim, M., Song, Y.: Homomorphic encryption for arithmetic of approximate numbers. In: Takagi, T., Peyrin, T. (eds.) ASIACRYPT 2017. LNCS, vol. 10624, pp. 409–437. Springer, Cham (2017). [https://doi.org/10.1007/978-3-319-70694-8\\_15](https://doi.org/10.1007/978-3-319-70694-8_15)
8. Chillotti, I., Gama, N., Georgieva, M., Izabachène, M.: TFHE: fast fully homomorphic encryption over the torus. J. Cryptol. **33**, 34–91 (2019)
9. Chillotti, I., Joye, M., Ligier, D., Orfila, J.B., Tap, S.: CONCRETE: concrete operates on ciphertexts rapidly by extending TFHE. In: WAHC 2020–8th Workshop on Encrypted Computing & Applied Homomorphic Cryptography, vol. 15 (2020)
10. Chillotti, I., Joye, M., Paillier, P.: Programmable bootstrapping enables efficient homomorphic inference of deep neural networks. IACR Cryptol. ePrint Arch. **2021**, 91 (2021)
11. Folkerts, L., Gouert, C., Tsoutsos, N.G.: REDsec: running encrypted discretized neural networks in seconds. In: Proceedings 2023 Network and Distributed System Security Symposium (2023)

---

<sup>1</sup> [https://github.com/zama-ai/concrete-ml/tree/release/0.6.x/use\\_case\\_examples](https://github.com/zama-ai/concrete-ml/tree/release/0.6.x/use_case_examples).

12. Gilad-Bachrach, R., Dowlin, N., Laine, K., Lauter, K., Naehrig, M., Wernsing, J.: CryptoNets: applying neural networks to encrypted data with high throughput and accuracy. In: Balcan, M.F., Weinberger, K.Q. (eds.) Proceedings of The 33rd International Conference on Machine Learning. Proceedings of Machine Learning Research, vol. 48, pp. 201–210. PMLR, New York (2016). <https://proceedings.mlr.press/v48/gilad-bachrach16.html>
13. Jacob, B., et al.: Quantization and training of neural networks for efficient integer-arithmetic-only inference. In: 2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition, pp. 2704–2713 (2017)
14. Joye, M.: Guide to fully homomorphic encryption over the [discretized] torus. Cryptology ePrint Archive, Paper 2021/1402 (2021). <https://eprint.iacr.org/2021/1402>, <https://eprint.iacr.org/2021/1402>
15. Lee, E., et al.: Low-complexity deep convolutional neural networks on fully homomorphic encryption using multiplexed parallel convolutions. In: Chaudhuri, K., Jegelka, S., Song, L., Szepesvari, C., Niu, G., Sabato, S. (eds.) Proceedings of the 39th International Conference on Machine Learning. Proceedings of Machine Learning Research, vol. 162, pp. 12403–12422. PMLR (2022). <https://proceedings.mlr.press/v162/lee22e.html>
16. Lou, Q., Jiang, L.: SHE: a fast and accurate deep neural network for encrypted data. In: Wallach, H., Larochelle, H., Beygelzimer, A., d’Alché-Buc, F., Fox, E., Garnett, R. (eds.) Advances in Neural Information Processing Systems, vol. 32. Curran Associates, Inc. (2019). <https://proceedings.neurips.cc/paper/2019/file/56a3107cad6611c8337ee36d178ca129-Paper.pdf>
17. Meyre, A., et al.: Concrete-ML (2022). <https://github.com/zama-ai/concrete-ml>
18. Pappalardo, A.: Xilinx/brevitas (2021). <https://doi.org/10.5281/zenodo.3333552>
19. Sejnowski, T.J.: The unreasonable effectiveness of deep learning in artificial intelligence. Proc. Natl. Acad. Sci. **117**(48), 30033–30038 (2020). <https://doi.org/10.1073/pnas.1907373117>, <https://www.pnas.org/doi/abs/10.1073/pnas.1907373117>
20. Yvinec, E., Dapogny, A., Cord, M., Bailly, K.: SPIQ: data-free per-channel static input quantization. CoRR abs/2203.14642 (2022). <https://doi.org/10.48550/arXiv.2203.14642>